Homework 1

Due on 03/03/2018

In this exercise, we will predict solubility of compounds using their chemical structures. The training data are in the file "solubility_train.csv" (951 observations) and the test data are in "solubility_test.csv" (356 observations). Among the 228 predictors, 208 are binary variables that indicate the presence or absence of a particular chemical substructure, 16 are count descriptors, such as the number of bonds or the number of bromine atoms, and 4 are continuous descriptors, such as molecular weight or surface area. The response is in the column "Solubility" (the last column).

(a) Fit a linear model using least squares on the training data and calculate the mean square error using the test data.

(b) Fit a ridge regression model on the training data, with λ chosen by cross-validation. Report the test error.

(c) Fit a lasso model on the training data, with λ chosen by cross-validation. Report the test error, along with the number of non-zero coefficient estimates.

(d) Fit a PCR model on the training data, with M chosen by cross-validation. Report the test error, along with the value of M selected by cross-validation.

(e) Briefly discuss the results obtained in (a)$\sim$(d).